

Problem II: K-means Clustering

1452669, Yang LI, April 7

Data Preprocessing

Since the raw trade data is high dimensional, I use Principal Component Analysis (PCA) to reduce them to a few principal components, in particular, I use 2 compared with the LSH method using Euclidean Distance. Hence it improve the average silhouette obviously.

K-means

K-means is the basic and comprehensive clustering algorithm, it accepts an initial set of k means $m_1^{(1)}, m_2^{(1)}, \dots, m_k^{(1)}$, and proceeds by alternating between two steps:

- Assignment: Assign each observation to the cluster whose mean has the least squared Euclidean distance.

$$S_i^{(t)} = \{x_p : \|x_p - m_i^{(t)}\|^2 \leq \|x_p - m_j^{(t)}\|^2 \forall j, 1 \leq j \leq k\}$$

where each x_p is assigned to exactly one $S_i^{(t)}$, even if it could be assigned to two or more of them.

- Update: Calculate the new means to be the centroids of the observations in the new clusters.

$$m_i^{(t+1)} = \frac{1}{|S_i^{(t)}|} \sum_{x_j \in S_i^{(t)}} x_j$$

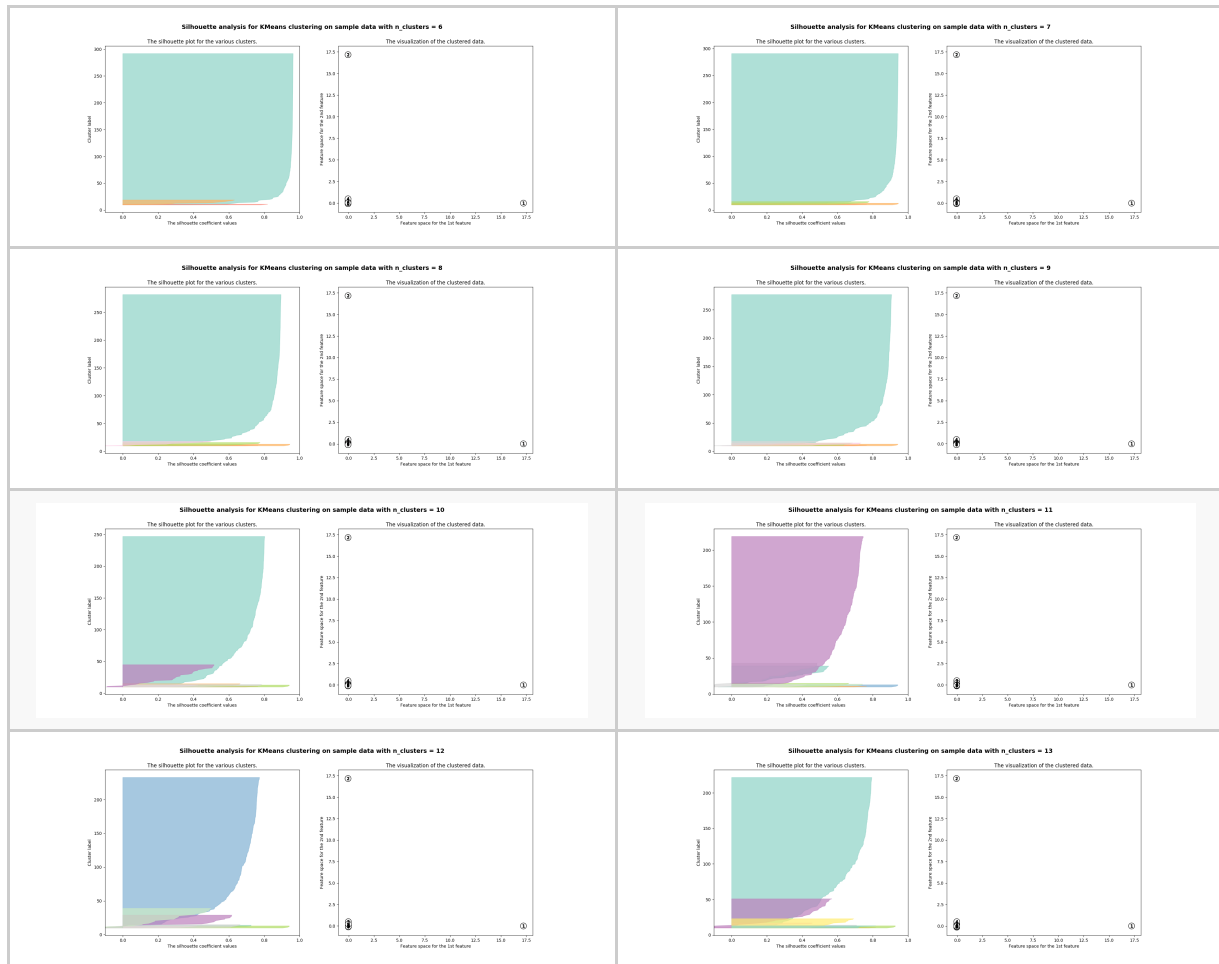
```
1 def kmeans(df, random_vip, knns):
2     k = int(math.sqrt(df.shape[1] / 2))
3     silhouette_avgs = []
4
5     for n_clusters in range(2, k + 2):
6         logging.debug("K-means: n_clusters = {}".format(n_clusters))
7         clusterer = KMeans(n_clusters=n_clusters)
8         X = PCA(n_components=2, whiten=True).fit_transform(df.T)
9         cluster_labels = clusterer.fit_predict(X)
10        silhouette_avg = silhouette_score(X, cluster_labels)
11        silhouette_avgs.append(silhouette_avg)
12        print("For n_clusters =", n_clusters,
13              "The average silhouette_score in K-means is :",
14              silhouette_avg)
15
16        # if n_clusters >= k / 2:
17        #     plot_silhouette(X, cluster_labels, n_clusters, clusterer)
18
19        res = 0
20        no = cluster_labels[df.columns.get_loc(random_vip)]
21        for neighbor in knns:
```

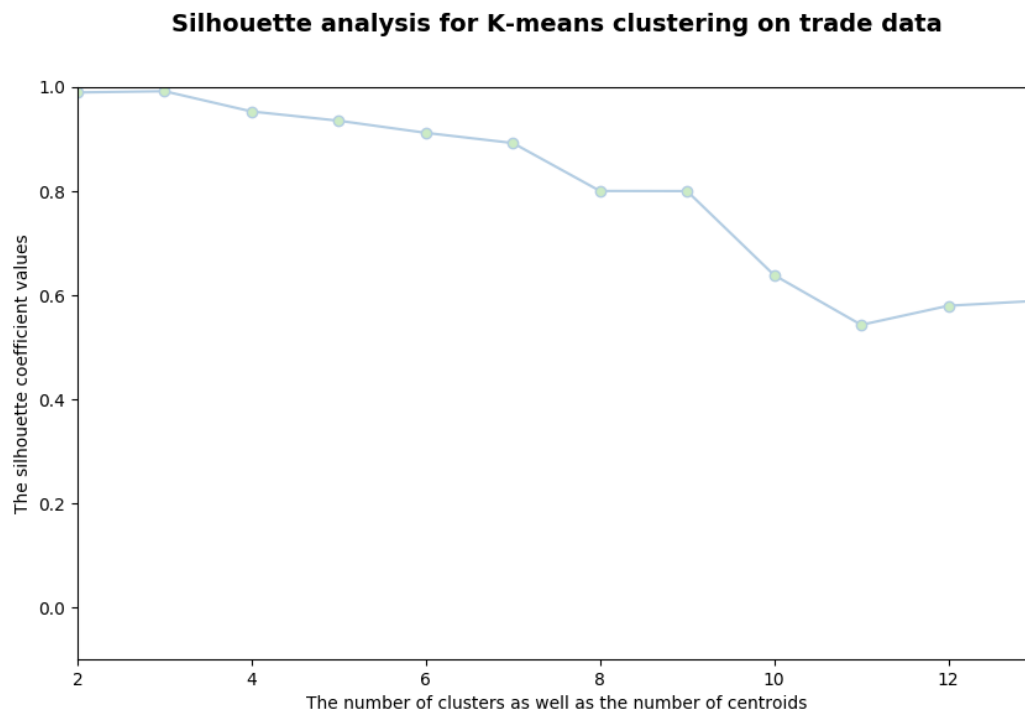
```

21         if cluster_labels[df.columns.get_loc(neighbor)] == no:
22             res += 1
23         else:
24             logging.info(
25                 "K-means: vipno: {} is not in the same
cluster.".format(
26                     neighbor))
27             print(
28                 "For k = {} in kNN, there has {} in the same cluster in K-
means.".format(
29                     len(knns), res))
30
31             # plot_kmeans_clusterno(k, silhouette_avgs)
32
33             return silhouette_avgs.index(max(silhouette_avgs)) + 2

```

The silhouette analysis results for different clusters show as follows:





Above figures intuitively indicate the number of clusters in the trade dataset, like 9 for example. The exact silhouette score are in the following table.

number of clusters	silhouette score
2	0.9898142095520185
3	0.9921749295434477
4	0.9532367585544059
5	0.9354793995672006
6	0.912177772818011
7	0.8925583120466012
8	0.8003854540541915
9	0.8001215931883985
10	0.6280774320073931
11	0.5553540967764518
12	0.6433141132870174
13	0.6210252967296004

Compared with the clustering of kNN using LSH, the result are in the following table (duplicate 100 times):

k	correctness
2	1.0
3	0.99
4	0.98
5	0.948
6	0.938
7	0.938
8	0.928
9	0.926
10	0.798
11	0.79
12	0.75
13	0.766

Performance

Time & Space Complexity in Theory

Let t_{dist} be the time to calculate the distance between two objects, K stands for the number of clusters (centroids) and n stands for the number of objects. Thus, given bound number of iterations I .

- time complexity: $O(IKnt_{dist}) = O(IKnm)$ where m means for m -dimensional vectors.
- space complexity: $O((n + K)m)$

Benchmark in Practice

Using the common algorithm for K-means clustering, it is a polynomial even though finding the optimal solution to the k-means problem is NP-hard. As we analyzed above, the Lloyd's algorithm is not the global optimum, it always gets a local maximum which differs from the random chosen at the beginning.

In practice, given a bound number of iteration, it cost both in data preprocessing and the finding and calculation of centroids parts.

Timer unit: 1e-06 s

Total time: 0.753593 s

File: /Users/Yang/Developer/420235DataMining/hw1/q2/kmeans.py

Function: kmeans at line 11

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
11					@profile
12					def kmeans(df, random_vip, knns):
13	1	176.0	176.0	0.0	k = int(math.sqrt(df.shape[1] / 2))
14	1	2.0	2.0	0.0	silhouette_avgs = []
15					
16	13	16.0	1.2	0.0	for n_clusters in range(2, k + 2):
17	12	1981.0	165.1	0.3	logging.debug("K-means: n_clusters = {}".format(n_clusters))
18	12	132.0	11.0	0.0	clusterer = KMeans(n_clusters=n_clusters)
19	12	294449.0	24537.4	39.1	X = PCA(n_components=2, whiten=True).fit_transform(df.T)
20	12	414335.0	34527.9	55.0	cluster_labels = clusterer.fit_predict(X)
21	12	41378.0	3448.2	5.5	silhouette_avg = silhouette_score(X, cluster_labels)
22	12	30.0	2.5	0.0	silhouette_avgs.append(silhouette_avg)
23	12	14.0	1.2	0.0	print("For n_clusters =", n_clusters,
24	12	368.0	30.7	0.0	"The average silhouette_score in K-means is :", silhouette_avg)
25					
26					# if n_clusters >= k / 2:
27					# plot_silhouette(X, cluster_labels, n_clusters, clusterer)
28					
29	12	12.0	1.0	0.0	res = 0
30	12	208.0	17.3	0.0	no = cluster_labels[df.columns.get_loc(random_vip)]
31	72	66.0	0.9	0.0	for neighbor in knns:
32	60	217.0	3.6	0.0	if cluster_labels[df.columns.get_loc(neighbor)] == no:
33	60	51.0	0.8	0.0	res += 1
34					else:
35					logging.info(
36					"K-means: vipno: {} is not in the same cluster.".format(
37					neighbor))
38	12	9.0	0.8	0.0	print(
39	12	11.0	0.9	0.0	"For k = {} in kNN, there has {} in the same cluster in K-means.".format(
40	12	134.0	11.2	0.0	len(knns), res))
41					
42					# plot_kmeans_clusterno(k, silhouette_avgs)
43					
44	1	4.0	4.0	0.0	return silhouette_avgs.index(max(silhouette_avgs)) + 2

Line #	Mem usage	Increment	Line Contents
=====			
11	137.504 MiB	137.504 MiB	@profile
12			def kmeans(df, random_vip, knns):
13	137.504 MiB	0.000 MiB	k = int(math.sqrt(df.shape[1] / 2))
14	137.504 MiB	0.000 MiB	silhouette_avgs = []
15			
16	147.027 MiB	0.000 MiB	for n_clusters in range(2, k + 2):
17	147.023 MiB	0.000 MiB	logging.debug("K-means: n_clusters = {}".format(n_clusters))
18	147.023 MiB	0.000 MiB	clusterer = KMeans(n_clusters=n_clusters)
19	147.027 MiB	8.465 MiB	X = PCA(n_components=2, whiten=True).fit_transform(df.T)
20	147.027 MiB	0.258 MiB	cluster_labels = clusterer.fit_predict(X)
21	147.027 MiB	0.781 MiB	silhouette_avg = silhouette_score(X, cluster_labels)
22	147.027 MiB	0.000 MiB	silhouette_avgs.append(silhouette_avg)
23	147.027 MiB	0.000 MiB	print("For n_clusters =", n_clusters,
24	147.027 MiB	0.020 MiB	"The average silhouette_score in K-means is :", silhouette_avg)
25			
26			# if n_clusters >= k / 2:
27			# plot_silhouette(X, cluster_labels, n_clusters, clusterer)
28			
29	147.027 MiB	0.000 MiB	res = 0
30	147.027 MiB	0.000 MiB	no = cluster_labels[df.columns.get_loc(random_vip)]
31	147.027 MiB	0.000 MiB	for neighbor in knns:
32	147.027 MiB	0.000 MiB	if cluster_labels[df.columns.get_loc(neighbor)] == no:
33	147.027 MiB	0.000 MiB	res += 1
34			else:
35			logging.info(
36			"K-means: vipno: {} is not in the same cluster.".format(
37			neighbor))
38	147.027 MiB	0.000 MiB	print(
39	147.027 MiB	0.000 MiB	"For k = {} in kNN, there has {} in the same cluster in K-means.".format(
40	147.027 MiB	0.000 MiB	len(knns), res))
41			
42			# plot_kmeans_clusterno(k, silhouette_avgs)
43			
44	147.027 MiB	0.000 MiB	return silhouette_avgs.index(max(silhouette_avgs)) + 2

Screenshot

```
/usr/local/bin/python3.6 /Users/Yang/Developer/420235DataMining/hw1/main.py
INFO:root:DataFrame shape: (2635, 298)
<class 'pandas.core.frame.DataFrame'>
Index: 2635 entries, 10000004 to 40000700
Columns: 298 entries, 13205496418 to 6222021615015662822
dtypes: float64(298)
memory usage: 6.0+ MB
DEBUG:root:DataFrame info: None
INFO:root:random vipno: 1590140606433
vipno in ranked order using kNN(k = 5):
1590160754770
1591011326672
1594140467704
1592140611301
1591020667889
DEBUG:root:K-means: n_clusters = 2
DEBUG:root:K-means: n_clusters = 3
For n_clusters = 2 The average silhouette_score in K-means is : 0.9898142095574196
For k = 5 in kNN, there has 5 in the same cluster in K-means.
DEBUG:root:K-means: n_clusters = 4
For n_clusters = 3 The average silhouette_score in K-means is : 0.9921749295445973
For k = 5 in kNN, there has 5 in the same cluster in K-means.
For n_clusters = 4 The average silhouette_score in K-means is : 0.9532367586043914
For k = 5 in kNN, there has 5 in the same cluster in K-means.
DEBUG:root:K-means: n_clusters = 5
For n_clusters = 5 The average silhouette_score in K-means is : 0.9354793989646298
For k = 5 in kNN, there has 5 in the same cluster in K-means.
DEBUG:root:K-means: n_clusters = 6
For n_clusters = 6 The average silhouette_score in K-means is : 0.9121777730153087
For k = 5 in kNN, there has 5 in the same cluster in K-means.
DEBUG:root:K-means: n_clusters = 7
DEBUG:root:K-means: n_clusters = 8
For n_clusters = 7 The average silhouette_score in K-means is : 0.8925583105475965
For k = 5 in kNN, there has 5 in the same cluster in K-means.
For n_clusters = 8 The average silhouette_score in K-means is : 0.8003854367303654
For k = 5 in kNN, there has 5 in the same cluster in K-means.
DEBUG:root:K-means: n_clusters = 9
DEBUG:root:K-means: n_clusters = 10
For n_clusters = 9 The average silhouette_score in K-means is : 0.8054482598322832
For k = 5 in kNN, there has 5 in the same cluster in K-means.
For n_clusters = 10 The average silhouette_score in K-means is : 0.6352370948494438
DEBUG:root:K-means: n_clusters = 11
For k = 5 in kNN, there has 5 in the same cluster in K-means.
For n_clusters = 11 The average silhouette_score in K-means is : 0.6411969201951703
DEBUG:root:K-means: n_clusters = 12
For k = 5 in kNN, there has 5 in the same cluster in K-means.
For n_clusters = 12 The average silhouette_score in K-means is : 0.636311067583867
For k = 5 in kNN, there has 5 in the same cluster in K-means.
DEBUG:root:K-means: n_clusters = 13
For n_clusters = 13 The average silhouette_score in K-means is : 0.6167497931869516
For k = 5 in kNN, there has 5 in the same cluster in K-means.

Process finished with exit code 0
```

